

Abstract

This paper presents the design and implementation of a comprehensive big data pipeline leveraging Amazon Web Services (AWS) for data storage and visualization, complemented by a Linux-based PySpark environment for distributed data processing. The implementation utilizes a substantial Flipkart e-commerce dataset containing over 5.7 million records, demonstrating the pipeline's capability to handle large-scale data processing requirements in a real-world context.

Introduction

E-commerce platforms generate massive amounts of data that require sophisticated processing and analysis capabilities. This project addresses the challenges of handling such large-scale data by implementing a robust big data pipeline. The solution combines the scalability of AWS services with the distributed processing power of PySpark, creating a comprehensive system for e-commerce data analysis.

The project utilizes a pre-crawled Flipkart dataset found on Kaggle, representing one of India's largest e-commerce platforms. The dataset encompasses:

- Size: 8.75 MB
- Records: 5.7+ million entries
- Format: Structured data with clear categorical and numerical attributes

The dataset contains crucial e-commerce metrics and attributes:

- **Product Identification**
 - `_id`: Unique identifier assigned to each product
 - `brand`: Product manufacturer or brand name
 - `category/sub_category`: Hierarchical product classification
- **Pricing Information**
 - `actual_price`: Original product price
 - `selling_price`: Discounted retail price
 - `discount`: Percentage reduction from original price
- **Performance Metrics**
 - `average_rating`: Aggregate customer satisfaction score
 - `out_of_stock`: Inventory availability indicator
 - `seller`: Vendor identification

This paper details the technical implementation, challenges encountered, and solutions developed throughout the project lifecycle, providing insights into building robust big data solutions for e-commerce applications.

Methodology

The initial phase involved setting up the AWS infrastructure for data storage:

- Created an S3 bucket for storing both raw and processed data
- Manually uploaded the raw dataset to the designated S3 bucket
- Established an Amazon Linux EC2 instance for processing
- Configured secure access using .pem to ppk conversion via PuttyGen
- Connected to EC2 instance using public IPv4 address through Putty

The screenshot shows the 'Create bucket' page in the AWS Management Console. The breadcrumb navigation is 'Amazon S3 > Buckets > Create bucket'. The page title is 'Create bucket' with an 'Info' link. Below the title, it states 'Buckets are containers for data stored in S3.' The 'General configuration' section includes the 'AWS Region' set to 'Europe (Stockholm) eu-north-1'. Under 'Bucket type', the 'General purpose' option is selected, with a description: 'Recommended for most use cases and access patterns. General purpose buckets are the original S3 bucket type. They allow a mix of storage classes that redundantly store objects across multiple Availability Zones.' The 'Directory' option is also visible, described as 'Recommended for low-latency use cases. These buckets use only the S3 Express One Zone storage class, which provides faster processing of data within a single Availability Zone.' The 'Bucket name' field contains 'anoushka-bda'. Below the field, a note states: 'Bucket name must be unique within the global namespace and follow the bucket naming rules. See rules for bucket naming'. Under 'Copy settings from existing bucket - optional', there is a 'Choose bucket' button and a format example: 'Format: s3://bucket/prefix'.

The screenshot shows the 'Buckets' page in the AWS Management Console. A green success message at the top states: 'Successfully created bucket "anoushka-bda"'. Below this, there is a section for 'Account snapshot - updated every 24 hours' with a 'View Storage Lens dashboard' button. The main content area is divided into 'General purpose buckets' and 'Directory buckets'. Under 'General purpose buckets (1)', there is a 'Create bucket' button and a table listing the bucket. The table has columns for 'Name', 'AWS Region', 'IAM Access Analyzer', and 'Creation date'. The listed bucket is 'anoushka-bda' in the 'Europe (Stockholm) eu-north-1' region, with a creation date of 'December 15, 2024, 10:07:01 (UTC-05:00)'.

Name	AWS Region	IAM Access Analyzer	Creation date
anoushka-bda	Europe (Stockholm) eu-north-1	View analyzer for eu-north-1	December 15, 2024, 10:07:01 (UTC-05:00)

The screenshot shows the AWS Management Console interface for the EC2 Instances page. The top navigation bar includes the AWS logo, a search bar, and various utility icons. The left sidebar contains navigation links for Dashboard, EC2 Global View, Events, and Instances. The main content area is titled 'Instances (1)' and includes a search bar, a 'Launch instances' button, and a table of instances. The table has columns for Name, Instance ID, Instance state, and Instance type. One instance, 'AnoushkaLinux', is listed with ID 'i-0f8e6c44b41b47d76' and is in the 'Running' state. Below the table, there is a 'Select an instance' section with a search bar and a dropdown menu.

Mini Project x Instance detail x EC2 Instance C... which java ver... How to install x +

eu-north-1.console.aws.amazon.com/ec2-instance-connect/ssh?region=eu-north-1&con...

FA24: SCIENTIFIC VI... Jobs applying - Goo... Top 2025 U.S. Entry... NHANES Tutorials ... Anoushka Shinde ... All Bookmarks

aws Search [Alt+S]

Amazon Linux 2023

<https://aws.amazon.com/linux/amazon-linux-2023>

Last login: Sun Dec 15 15:46:40 2024 from 45.119.30.69

[ec2-user@ip-172-31-27-220 ~]\$

i-0f8e6c44b41b47d76 (AnoushkaLinux)

PublicIPs: 13.48.135.5 PrivateIPs: 172.31.27.220

```

Last login: Sun Dec 15 15:21:13 2024 from 13.48.4.203
[ec2-user@ip-172-31-27-220 ~]$ uname -a
Linux ip-172-31-27-220.eu-north-1.compute.internal 6.1.115-126.197.amzn2023.x86_
64 #1 SMP PREEMPT_DYNAMIC Tue Nov 5 17:36:57 UTC 2024 x86_64 x86_64 x86_64 GNU/
Linux
[ec2-user@ip-172-31-27-220 ~]$ df -h
Filesystem      Size  Used Avail Use% Mounted on
devtmpfs        4.0M   0  4.0M   0% /dev
tmpfs           453M   0  453M   0% /dev/shm
tmpfs           181M  436K  181M   1% /run
/dev/nvme0n1p1   30G   1.7G   29G   6% /
tmpfs           453M   0  453M   0% /tmp
/dev/nvme0n1p128 10M   1.3M   8.7M  13% /boot/efi
tmpfs           91M   0   91M   0% /run/user/1000
[ec2-user@ip-172-31-27-220 ~]$

```

The PySpark environment setup encountered and resolved several technical challenges:

```
rm -rf ~/temp_install
```

```
mkdir -p ~/temp_install
```

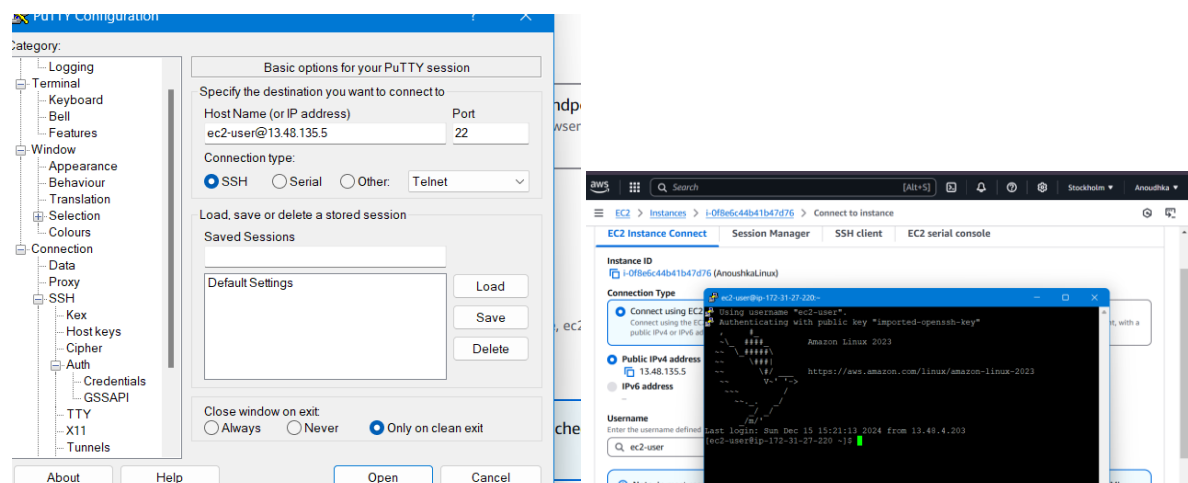
```
chmod 777 ~/temp_install
```

```
TMPDIR=~/temp_install pip3 install --no-cache-dir 'pyspark==3.2.0'
```

This configuration addressed initial space constraints by:

- Cleaning up the temporary directory
- Creating a new directory with explicit permissions
- Installing a specific PySpark version with temporary directory specification

I also set up Putty for accessing it through my desktop. I did this by converting the .pem file to a .ppk file using PuttyGen and then accessing it through Putty.



Another issue I faced is that I tried to configure aws using aws configure command. And put in random values. Then I forgot about it and created an IAM role. I momentarily forgot that I had listed values and so to fix this I put –

aws configure list

rm -rf ~/.aws/credentials ~/.aws/config

The screenshot shows the AWS IAM console. The top navigation bar includes the AWS logo, a search bar, and the user's name 'Anoushka'. The left sidebar shows the 'Identity and Access Management (IAM)' menu with options like Dashboard, Access management, and Access reports. The main content area displays the details for the 'anoushkacli' role, including its creation date, ARN, and instance profile ARN. Below this, the 'Permissions' tab is selected, showing a list of permissions policies. The bottom section shows the 'Modify IAM role' page, where the 'Instance ID' is 'i-0f8e6c44b41b47d76' and the 'IAM role' is 'anoushkacli'. The 'Update IAM role' button is highlighted in orange.

I then verified the s3 access from ec2.

```
[ec2-user@ip-172-31-27-220 ~]$ rm -rf ~/.aws/credentials ~/.aws/config
[ec2-user@ip-172-31-27-220 ~]$ aws s3 ls
2024-12-15 15:07:03 anoushka-bda
```

Run the following AWS CLI command to list the bucket:

```
[ec2-user@ip-172-31-27-220 ~]$ aws s3 ls s3://anoushka-bda/
2024-12-15 18:45:42 50725653 output.csv
[ec2-user@ip-172-31-27-220 ~]$
```

PySpark was initialized with necessary AWS support:

```
pyspark --packages org.apache.hadoop:hadoop-aws:3.3.1
```

```
[ec2-user@ip-172-31-27-220 ~]$ pyspark --packages org.apache.hadoop:hadoop-aws:3.3.1
Python 3.9.16 (main, Jul 5 2024, 00:00:00)
[GCC 11.4.1 20230605 (Red Hat 11.4.1-2)] on linux
```

And then I inspected the dataset for common information.

```
>>> df = spark.read.csv("s3a://anoushka-bda/output.csv",header = True, inferSchema = True)
24/12/16 07:18:50 WARN MetricsConfig: Cannot locate configuration: tried hadoop-metrics2-s3a-file-system.properties,hadoop-metrics2.properties
>>> # Display the first 5 rows
>>> print("Sample Data from the Dataset:")
Sample Data from the Dataset:
>>> df.show(5)
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|Unnamed: 0|      pid|      id|actual_price|average_rating|brand|      category|      crawled_at|description|      discount|      images|out_
of_stock|product_details|seller|selling_price|sub_category|title|url|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|0|fa8e22d6-c0b6-522...|      2,999|      3.9|York|Clothing and Acce...|02/10/2021, 20:11:51|      69% off|Yorker trackpants...|['https://rukmini...
False|TKPFCZ9EA7H5FYZH|['Style Code': '...|Shyam Enterprises|      921|      Bottomwear|Solid Men Multico...|https://www.flipk...
|1|893e6980-f2a0-531...|      1,499|      3.9|York|Clothing and Acce...|02/10/2021, 20:11:52|      66% off|Yorker trackpants...|['https://rukmini...
False|TKPFCZ9E9VZUVRZ|['Style Code': '...|Shyam Enterprises|      499|      Bottomwear|Solid Men Blue Tr...|https://www.flipk...
|2|eb4c8eab-8206-59d...|      2,999|      3.9|York|Clothing and Acce...|02/10/2021, 20:11:52|      68% off|Yorker trackpants...|['https://rukmini...
False|TKPFCZ9E9FV3Z4V|['Style Code': '...|Shyam Enterprises|      931|      Bottomwear|Solid Men Multico...|https://www.flipk...
|3|3f3f97bb-5faf-57d...|      2,999|      3.9|York|Clothing and Acce...|02/10/2021, 20:11:53|      69% off|Yorker trackpants...|['https://rukmini...
False|TKPFCZ9ESZ27YWEF|['Style Code': '...|Shyam Enterprises|      911|      Bottomwear|Solid Men Multico...|https://www.flipk...
|4|750caa3d-6264-53c...|      2,999|      3.9|York|Clothing and Acce...|02/10/2021, 20:11:53|      68% off|Yorker trackpants...|['https://rukmini...
False|TKPFCZ9EVXKBSUD7|['Style Code': '...|Shyam Enterprises|      943|      Bottomwear|Solid Men Brown, ...|https://www.flipk...
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
```

```
>>> df.printSchema()
root
 |-- Unnamed: 0: string (nullable = true)
 |-- _id: string (nullable = true)
 |-- actual_price: string (nullable = true)
 |-- average_rating: string (nullable = true)
 |-- brand: string (nullable = true)
 |-- category: string (nullable = true)
 |-- crawled_at: string (nullable = true)
 |-- description: string (nullable = true)
 |-- discount: string (nullable = true)
 |-- images: string (nullable = true)
 |-- out_of_stock: string (nullable = true)
 |-- pid: string (nullable = true)
 |-- product_details: string (nullable = true)
 |-- seller: string (nullable = true)
 |-- selling_price: string (nullable = true)
 |-- sub_category: string (nullable = true)
 |-- title: string (nullable = true)
 |-- url: string (nullable = true)

>>> #Total Rows
>>> print("Total Rows in the Dataset:", df.count())
[Stage 3:> (0 + 2) / 2]
Total Rows in the Dataset: 31518
>>>
>>> |
```

i-0f8e6c44b41b47d76 (AnoushkaLinux)

The data cleaning process involved several key steps:

1. Null Value Management
 - Implemented systematic null value removal
2. Data Type Conversion
 - Converted actual_price and selling_price to numeric format
 - Standardized discount column by removing 'off' suffix
3. Column Standardization
 - Corrected column name mismatches between description and discount
 - Implemented proper column renaming

4. Feature Engineering New columns created:

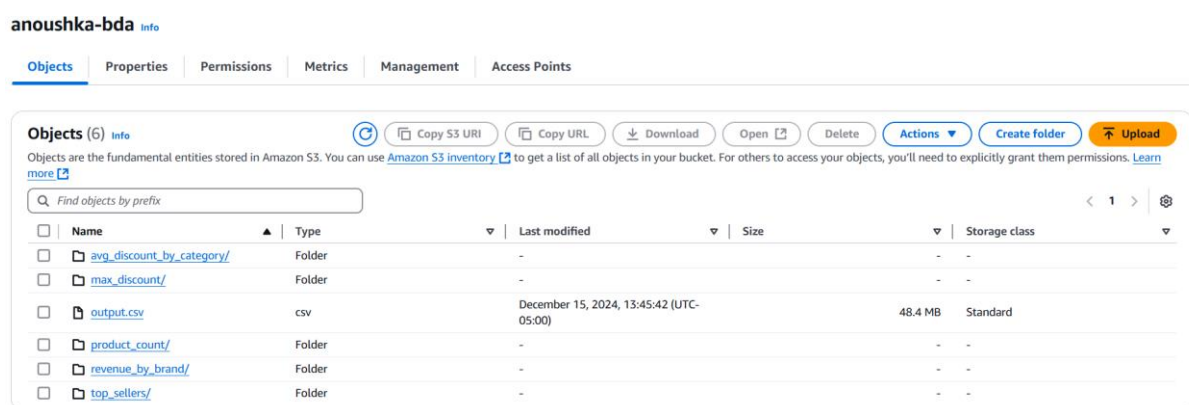
- Crawled_time: Timestamp information
- discounted_price: How much price was discounted

```
>>> df_final.show(5)
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| pid|      id|actual_price|average_rating|brand|seller|selling_price|sub_category|      crawled_at|discount|description|      images|out_of_stock|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|fa8e22d6-c0b6-522...|2999.0|      3.9|York|Clothing and Acce...|02/10/2021, 20:11:51|69.0|Yorker trackpants...|['https://rukmini...|False|TKPPF
29EA7H5FY2H|(['Style Code': '...|Shyam Enterprises|921.0|Bottomwear|Solid Men Multico...|https://www.flipk...|20:11:51|2078.0|
|893e6980-f2a0-531...|1499.0|      3.9|York|Clothing and Acce...|02/10/2021, 20:11:52|66.0|Yorker trackpants...|['https://rukmini...|False|TKPPF
29EJZV2UVR2|(['Style Code': '...|Shyam Enterprises|499.0|Bottomwear|Solid Men Blue Tr...|https://www.flipk...|20:11:52|1000.0|
|eb4c8eab-8206-59d...|2999.0|      3.9|York|Clothing and Acce...|02/10/2021, 20:11:52|68.0|Yorker trackpants...|['https://rukmini...|False|TKPPF
29EHFCY5Z4Y|(['Style Code': '...|Shyam Enterprises|931.0|Bottomwear|Solid Men Multico...|https://www.flipk...|20:11:52|2068.0|
|3f3497b3-5faf-57d...|2999.0|      3.9|York|Clothing and Acce...|02/10/2021, 20:11:53|69.0|Yorker trackpants...|['https://rukmini...|False|TKPPF
29KSZZ7WWEF|(['Style Code': '...|Shyam Enterprises|911.0|Bottomwear|Solid Men Multico...|https://www.flipk...|20:11:53|2088.0|
|750caa3d-6264-53c...|2999.0|      3.9|York|Clothing and Acce...|02/10/2021, 20:11:53|68.0|Yorker trackpants...|['https://rukmini...|False|TKPPF
29EVXKBSUD7|(['Style Code': '...|Shyam Enterprises|943.0|Bottomwear|Solid Men Brown, ...|https://www.flipk...|20:11:53|2056.0|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 5 rows
>>>
```

i-0f8e6c44b41b47d76 (AnoushkaLinux)
PublicIPs: 51.20.98.100 PrivateIPs: 172.31.27.220

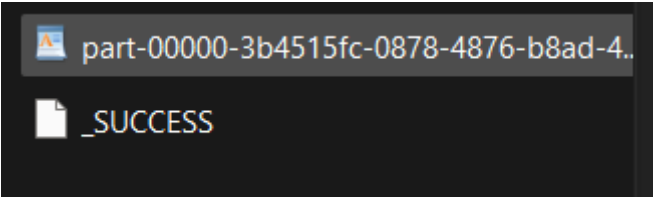
The project implemented two comprehensive aggregation approaches, with results stored directly in the S3 bucket:

1. Python-based Aggregations: The Python implementation focused on key business metrics and was executed through a dedicated .py script. The following aggregations were performed:
 - Revenue by brand: Calculated total revenue generation across different brands
 - Average discount by category: Analyzed discount patterns within product categories
 - Top sellers analysis: Identified and ranked the most successful sellers
 - Product count metrics: Quantified product distribution
 - Maximum discount price calculations: Determined highest discount offerings



Implementation Note: The output generation process created folders containing two files:

- A success indicator file
- A CSV file containing the aggregation results



Post-processing steps included:

- Manual download of generated files from S3
- Systematic renaming of CSV files to reflect their specific tasks
- Organization into a dedicated 'python' folder within the output directory
- Re-upload of the organized files to the S3 bucket

avg_discount_by_category/ Copy S3 URI

Objects Properties

Objects (2) Info

Copy S3 URI Copy URL Download Open Delete Actions Create folder Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	_SUCCESS	-	December 16, 2024, 12:55:59 (UTC-05:00)	0 B	Standard
<input type="checkbox"/>	avg_discount_by_category.csv	csv	December 16, 2024, 12:55:59 (UTC-05:00)	102.0 B	Standard

All the outputs can be found here.

Results –

Revenue by brand

brand	Total_Revenue
K	14736.0
fas	1495.0
ECKO Unl	522840.0
ertugrul ghazi c	5519.0
SATDEVANGIKHADIBHAND	73026.0
Marca Disa	122597.0
Aeload	749.0
A J STYL	23512.0
Jagdish Garmen	10541.0
NEXXE	8328.0
Gracew	153270.0
Pramukhraj Enterpri	185.0
CupidSto	160296.0
Musk	276.0
ATTIITU	75466.0
The Dry Ca	14838.0
DISCOUNT OUTL	9538.0
My lifestylebazz	798.0
BS S	769.0
Royalty Retail And Expo	6444.0
Man	2160.0
GLAUBEN PO	1895.0
EYETWIST	2457.0
style acco	3865.0
Vector	145229.0
LA SMI	234.0
U.S. POLO ASS	83391.0
She	349.0
Jack Roy	529.0
ANGI	280.0
Bone	47882.0
NEXT	1136.0
Shoef	218824.0
MARATH	1285.0
Fairdea	20707.0
Clo	13899.0

Avg discount by category

category	Avg_Discount_Percentage
Footwear	29.3271375464684
Clothing and Accessories	49.51085052920125

Top sellers

```
seller,Total_Sales
RetailNet,1420146.0
SandSMarketing,522840.0
Tayab Manch Fashions,406257.0
BioworldMerchandising,394641.0
ARBOR,302345.0
```

Product count

```
sub_category,Product_Count
Tracksuits,20
Inspire Clothing and Accessories,21
"Blazers, Waistcoats and Suits",26
INSPIRE Clothing and Accessories,4
Clothing Accessories,739
Men's Footwear,269
Topwear,8381
"Kurtas, Ethnic Sets and Bottoms",361
Sleepwear,42
Crocks Club Clothing and Accessories,10
Winter Wear,1259
Raincoats,25
Bottomwear,1753
Innerwear and Swimwear,443
Fabrics,22
Uber Urban Clothing and Accessories,27
```

Max discount price

```
Max_Discount_Percentage
7800.0
```

2. SQL-based Aggregations: SQL aggregations were implemented through PySpark's SQL interface, providing different analytical perspectives:

- Top discounted products: Identification of products with highest discount percentages
- Average rating per brand: Brand performance analysis based on customer ratings
- Product count per category: Detailed category-wise inventory analysis
- Most popular seller analysis: Seller performance metrics and rankings
- Out-of-stock products: Inventory status tracking and analysis

The SQL implementation was documented in a separate .py file, ensuring reproducibility and maintainability of the analysis pipeline. Each aggregation was designed to provide specific business insights while maintaining efficient processing of the large-scale dataset.

Similarly, I completed the code for sql in pyspark. The entire code has been pasted in a .py file that is present in the attachments.

Destination s3://anoushka-bda/sqlquery/	Succeeded 5 files, 229.6 KB (100.00%)	Failed 0 files, 0 B (0%)
--	--	-----------------------------

Files and folders	Configuration
-------------------	---------------

Files and folders (5 total, 229.6 KB)

< 1 >

Name	Folder	Type	Size	Status	Error
most_popular_seller.csv	-	text/csv	27.0 B	Succeeded	-
top_discounted_products.csv	-	text/csv	2.3 KB	Succeeded	-
product_count_per_category.csv	-	text/csv	132.9 KB	Succeeded	-
out_of_stock_products.csv	-	text/csv	27.0 B	Succeeded	-
avg_rating_per_brand.csv	-	text/csv	94.4 KB	Succeeded	-

Results –

Top Discounted products

```

_id,title,actual_price,selling_price,discount
86d0e25f-ddbe-5807-ac46-070ecd14db48,Printed Men Polo Neck Blue
T-ShirtÃ Å (Pack of 3),5500.0,950.0,
57b8f120-68e2-57b4-a961-12cbef8974d0,"Printed, Striped Men Round
Neck Blue T-Shirt",4999.0,800.0,
8ec86a28-1200-56c6-9dba-4c9e0f6d1fd9,Full Sleeve Solid Men Padded
Jacket,4999.0,995.0,
619f5e66-c5be-50e0-8bf5-a78d44b9b394,Striped Men V Neck
Multicolor T-Shirt,4200.0,799.0,
e8932882-c6ce-5f0d-b0ef-f330c5186e27,Solid Men Round Neck Light
Blue T-ShirtÃ Å (Pack of 5),3995.0,799.0,"Refresh your wardrobe
with the Premium Collection of Cotton Round Neck T-shirts from
Scott International. These t-shirts are made of Pure Cotton and
provide comfort all day. Melange Fabric, Elegant Stitching and
finish, Soft Neck Tape, Comfort Fit makes this product an ideal
daily wear and casual wear. Combine with Trousers or Denims for a
cool Casual look. Combo pack is ideal value for money."
65b654a6-1a14-55c5-b497-56eff5d4066a,"Solid Men Round Neck
Maroon, Blue, Light Blue, Dark Blue, Black T-ShirtÃ Å (Pack of
5)",3995.0,799.0,
9e466d9e-f3ef-5516-8f2e-3000b83b90bd,"Solid Men Round Neck
Maroon, Blue, Light Blue, Dark Blue, Black T-ShirtÃ Å (Pack of
5)",3995.0,799.0,
e926483d-80d3-5c8d-86c5-76ae4b3781f6,Solid Men Round Neck
Multicolor T-ShirtÃ Å (Pack of 5),3995.0,799.0,"Refresh your
wardrobe with the Premium Collection of Cotton Round Neck T-
shirts from Scott International. These t-shirts are made of Pure
Cotton and provide comfort all day. Melange Fabric, Elegant
Stitching and finish, Soft Neck Tape, Comfort Fit makes this
product an ideal daily wear and casual wear. Combine with
Trousers or Denims for a cool Casual look. Combo pack is ideal

```

Avg rating per brand

```

brand,avg_rating
62,61.0
MILD,5.0
VARTe,5.0
VIKING.INE,4.640000000000001
SORA,4.609459459459455
COL,4.55
Modest Ci,4.5
PixF,4.468922305764411
Cots,4.456357388316171
GYMBROTHe,4.407142857142857
Royalta,4.4
Ziko,4.4
JUARI BE A GENTLEM,4.4
Neon Ro,4.4
COOL CRU,4.4
LOUIS MONAR,4.3999999999999995
Bofri,4.366666666666667
Mega Sty,4.357142857142857
Zack Fo,4.354651162790697
Onei,4.302531645569618
Bracev,4.3
U.S. Polo Ass,4.3
ud fabr,4.3
Fresh Fe,4.299999999999999
Kt Kun,4.266666666666667
U.S.Polo As,4.23548387096774
Gracew,4.23118811881188
Aeload,4.225
limited colour,4.210169491525429
Truemo,4.204166666666668
ALQI,4.200000000000001
ANGI,4.2
MARATH,4.2
BOYT,4.2
Mohan Reta,4.2

```

Product count per category

```
category,product_count
Clothing and Accessories,28971
,1093
Footwear,987
"Bags, Wallets & Belts",41
False,35
T10Sports,34
EighteenUp,14
ATTITUDE,14
T T Limited,14
Tayab Manch Fashions,12
RetailNet,10
Assiduus Distribution,10
MILDIN(New Sell,7
PollyStores,6
DressDen,4
FABUNIFORMS,4
Length-19,4
XXL Sizes.,4
Buy More,3
ModaElementi,2
this much-awaited cricket season.,2
Suhani Garment,2
```

Most popular seller

```
|seller,product_count
,3422
```

Out of stock products

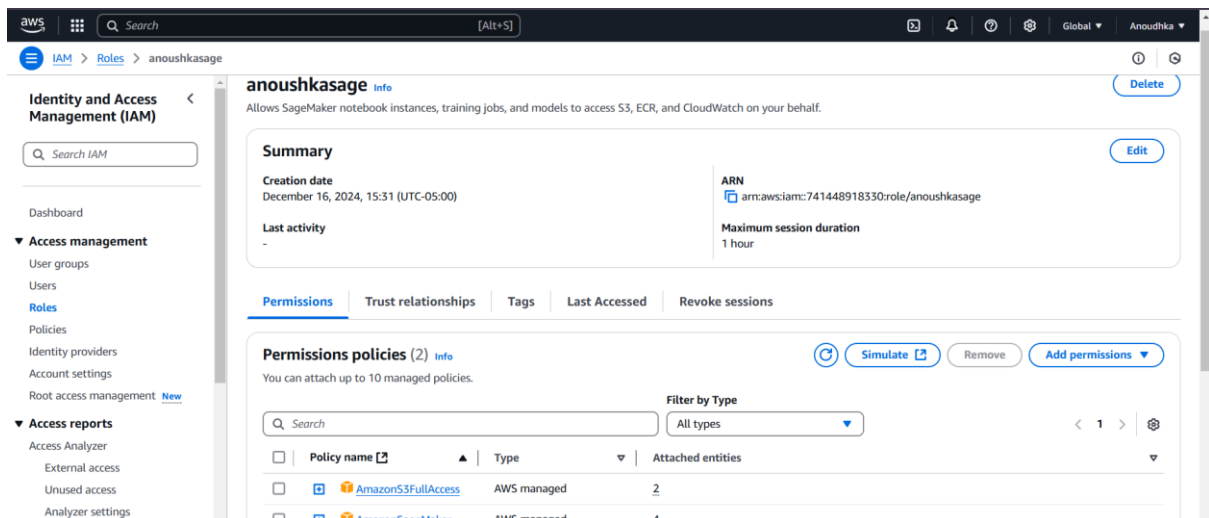
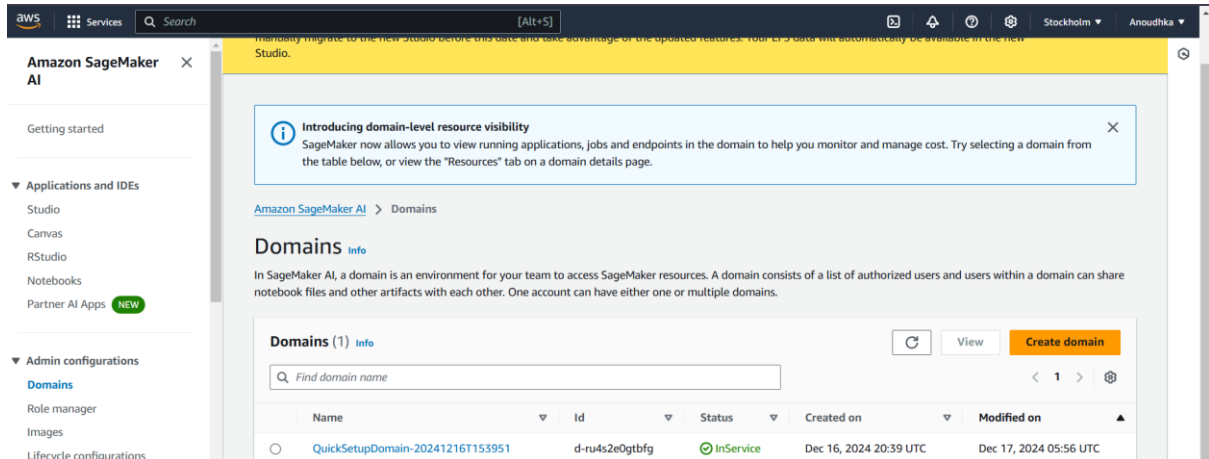
```
title,brand,category
499,,Clothing and Accessories
nu-lite Satin Tie & CufflinkÂ (Red),,Clothing and Accessories
so get this stunning necktie for your husband for your
anniversary,,Clothing and Accessories
so get this stunning necktie for your husband for your
anniversary,,Clothing and Accessories
Striped Collared Neck Casual Men Green Sweater,Man,Clothing and
Accessories
Solid V Neck Casual Men Grey Sweater,Man,Clothing and Accessories
Self Design Men Polo Neck Dark Blue T-Shirt,DISCOUNT
OUTL,Clothing and Accessories
Self Design Men Polo Neck Grey T-Shirt,DISCOUNT OUTL,Clothing and
Accessories
Self Design Men Polo Neck Blue T-Shirt,DISCOUNT OUTL,Clothing and
Accessories
Solid Men Polo Neck Red T-Shirt,DISCOUNT OUTL,Clothing and
Accessories
Self Design Men Polo Neck Beige T-Shirt,DISCOUNT OUTL,Clothing
and Accessories
,DISCOUNT OUTL,Clothing and Accessories
Men Tailored Fit Solid Casual Shirt,DISCOUNT OUTL,Clothing and
Accessories
Men Tailored Fit Solid Casual Shirt,DISCOUNT OUTL,Clothing and
Accessories
Men Tailored Fit Solid Casual Shirt,DISCOUNT OUTL,Clothing and
Accessories
Men Tailored Fit Solid Casual Shirt,DISCOUNT OUTL,Clothing and
Accessories
Printed Men Round Neck Orange T-Shirt,adidas Origina,Clothing and
Accessories
Solid Men Round Neck Black T-Shirt,adidas Origina,Clothing and
Accessories
Sporty Men Round Neck Black T-Shirt,adidas Origina,Clothing and
Accessories
Solid Men Dark Blue Sports Shorts,adidas Origina,Clothing and
Accessories
```

The project leveraged AWS SageMaker Autopilot for implementing machine learning capabilities, focusing on automated model training and evaluation without direct code intervention. This approach enabled efficient model development while maintaining high standards of accuracy and reliability. The initial setup phase involved several critical steps to ensure proper integration with AWS services:

1. Domain Configuration

- A new SageMaker domain was established to provide a controlled environment for machine learning operations

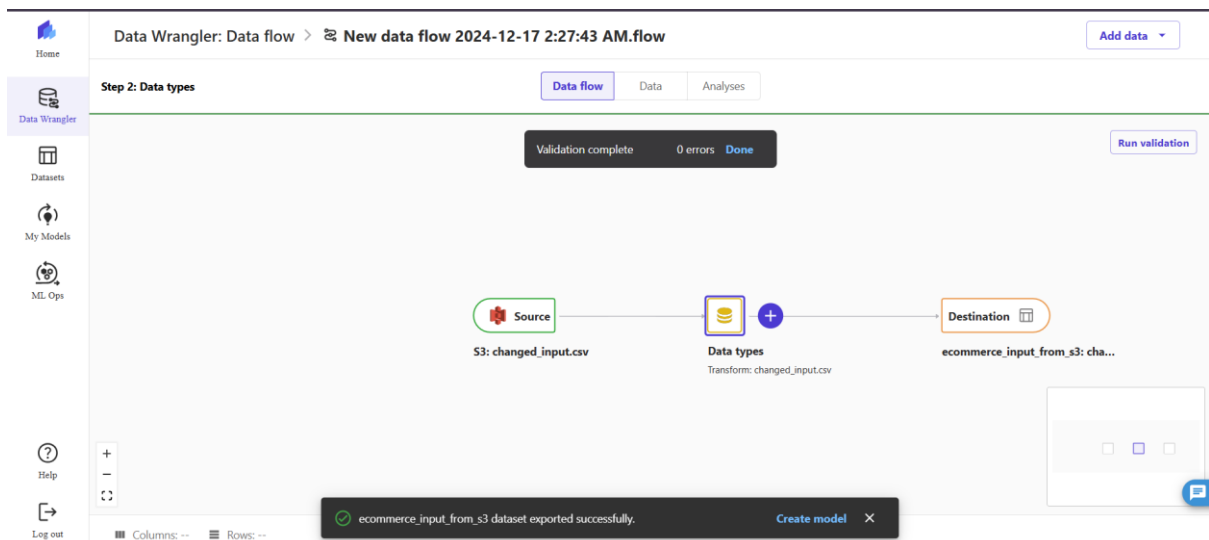
- The environment was configured with a specialized IAM role to manage necessary permissions and access controls
- Security protocols were implemented to ensure safe data handling and model training



The data integration process focused on seamlessly connecting the processed dataset with SageMaker Autopilot:

1. Data Import Process

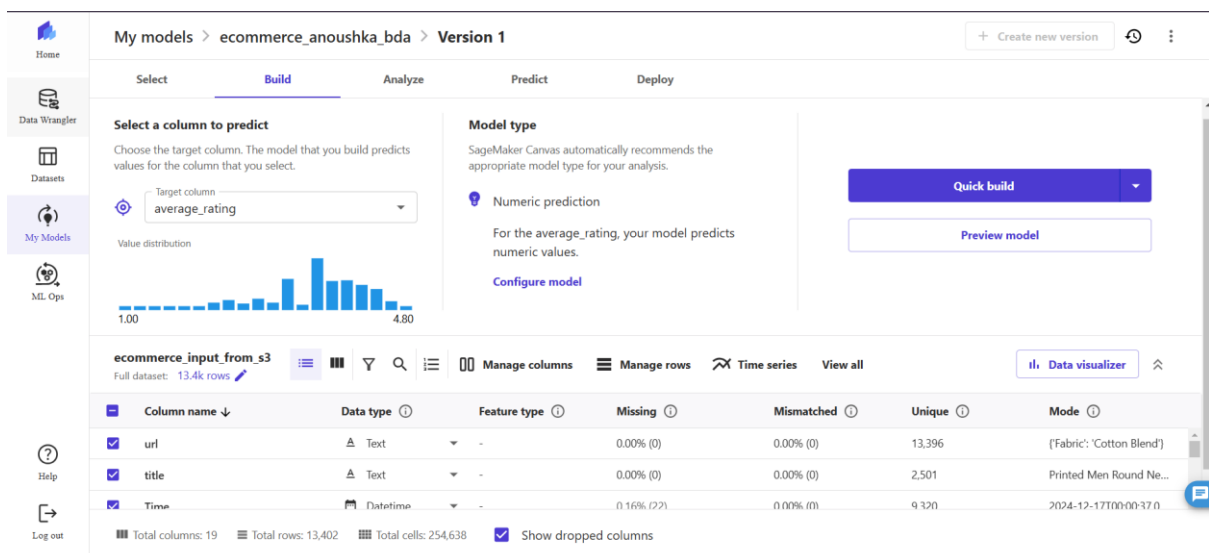
- Established a direct connection to the processed dataset stored in the S3 bucket
- Implemented automated data loading procedures to ensure efficient data transfer
- Verified data integrity throughout the import process



The automated machine learning process was executed through a structured approach:

1. Model Training Configuration

- Carefully selected appropriate target variables for prediction
- Configured the AutoML process to explore multiple model architectures
- Established evaluation criteria for model performance assessment



2. Training Execution

- Initiated the automated model training process
- Monitored the training progress across multiple model iterations
- Documented computational resource utilization and training duration

Select Build **Analyze** Predict Deploy

Model overview

Your model is being created. Standard build usually takes between 2-4 hours. You can now leave this view.

Time elapsed
8 min 17 sec

Expected build time
45 min

Build type
Standard build

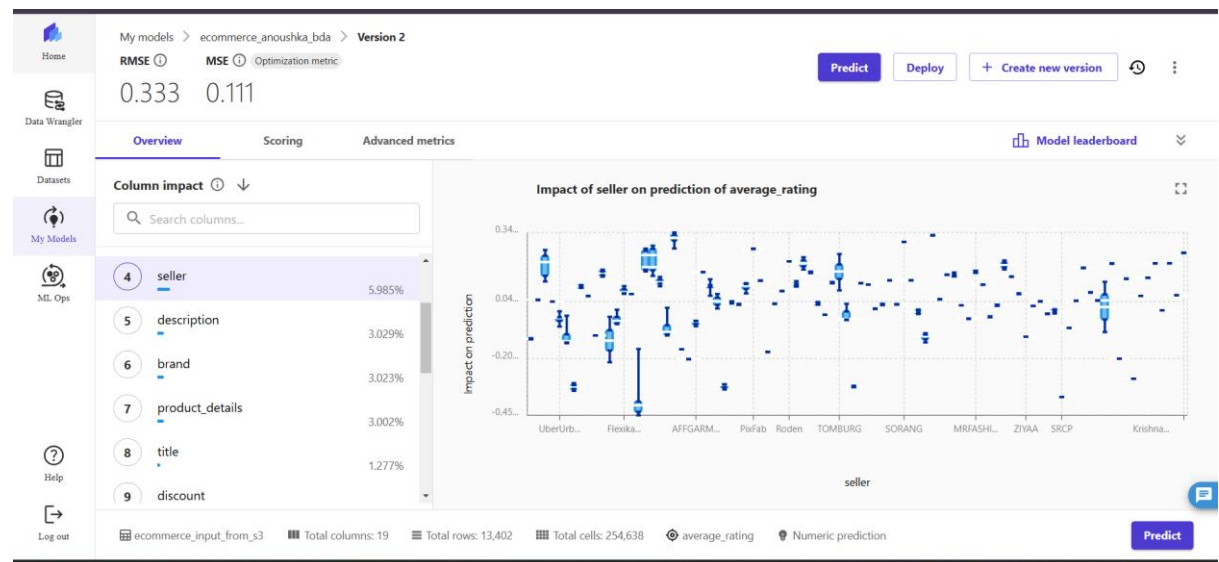
Detailed progress
Training models

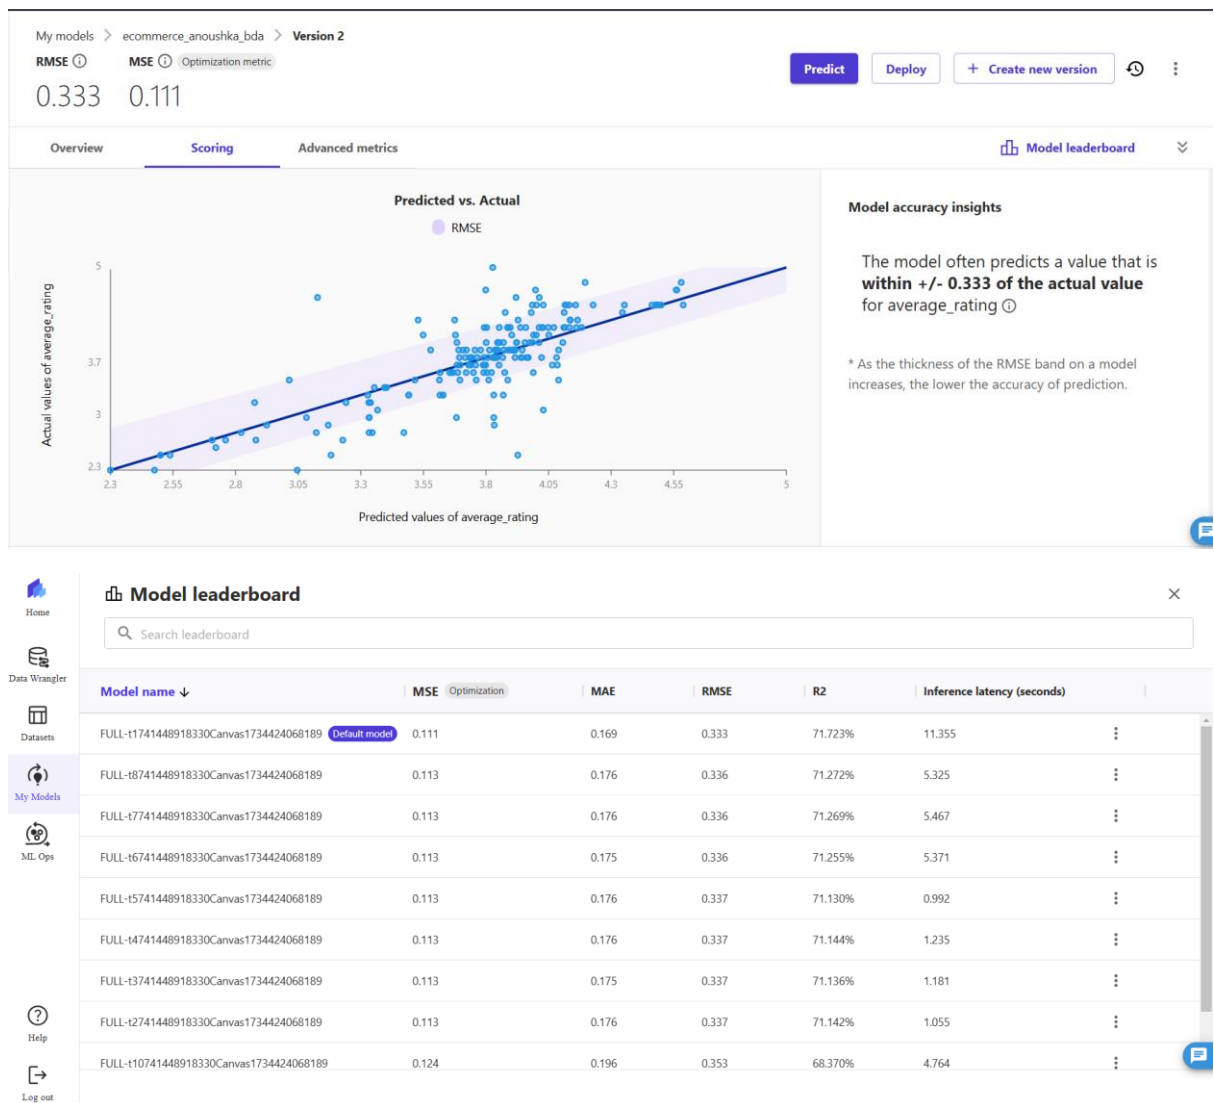


ecommerce_input_from_s3 Total columns: 19 Total rows: 13,402 Total cells: 254,638 average_rating Numeric prediction

3. Model Evaluation

- Analyzed the model leaderboard to identify top-performing algorithms
- Evaluated comprehensive performance metrics across all trained models
- Documented comparative analysis of different model architectures

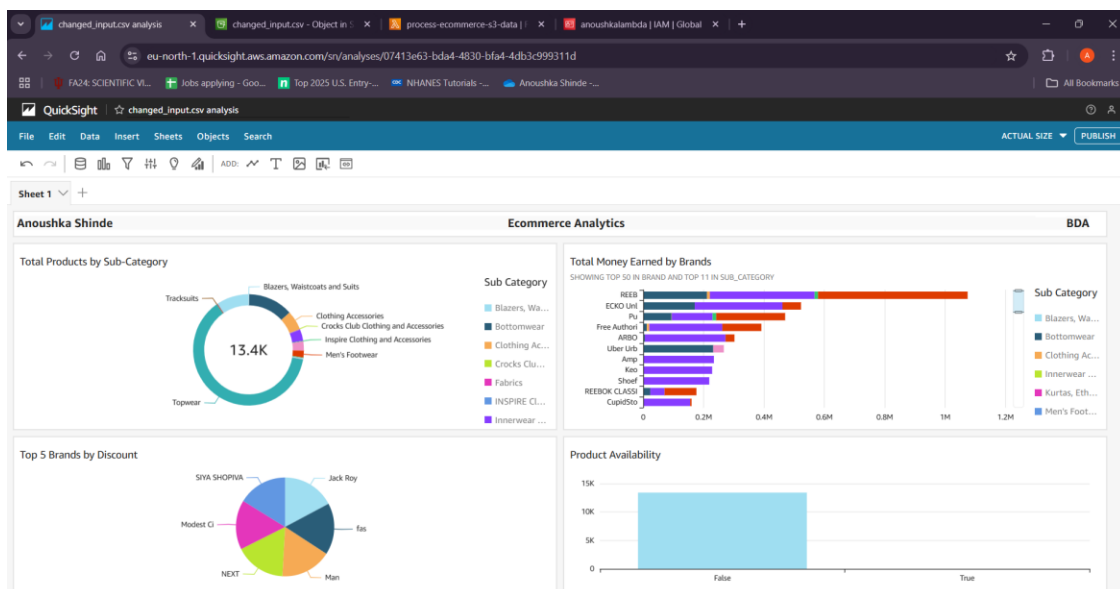
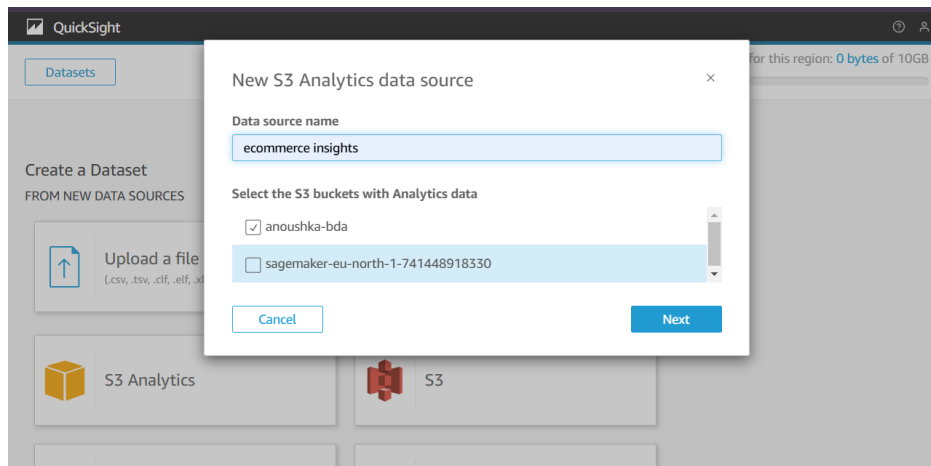




The best model out of these was the default model because it has :

- **Lowest MSE and RMSE:** Indicates the model has the least error in predictions.
- **Highest R^2 :** Shows the model explains the highest percentage of the variance in the data.

To visualise this data I created an account with Amazon Quicksight and gave access to my S3 bucket in the beginning of initialization itself. Then I created a new dataset where data was sourced from my S3 bucket.

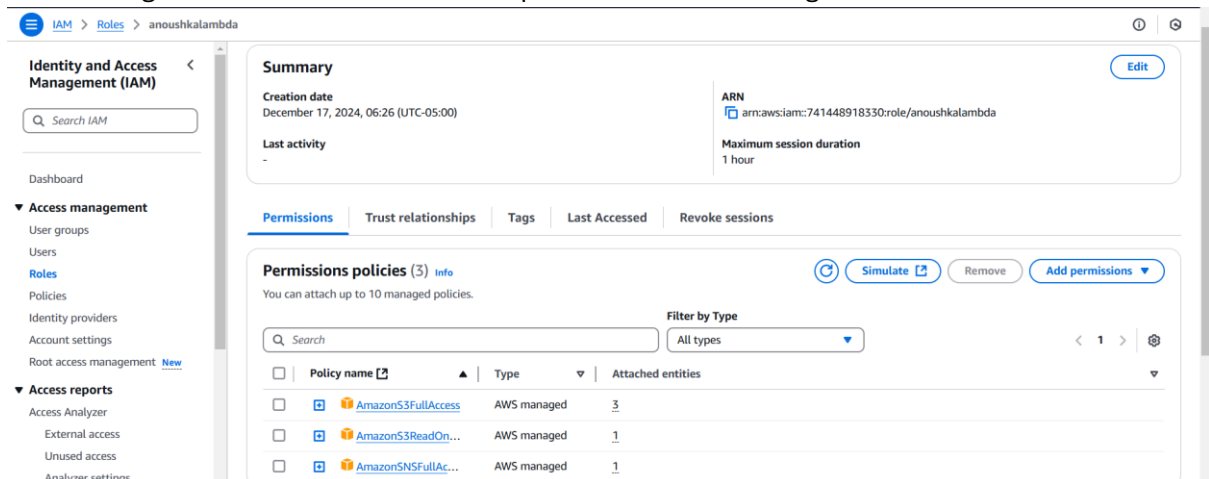


From the above visualizations we can summarise that:

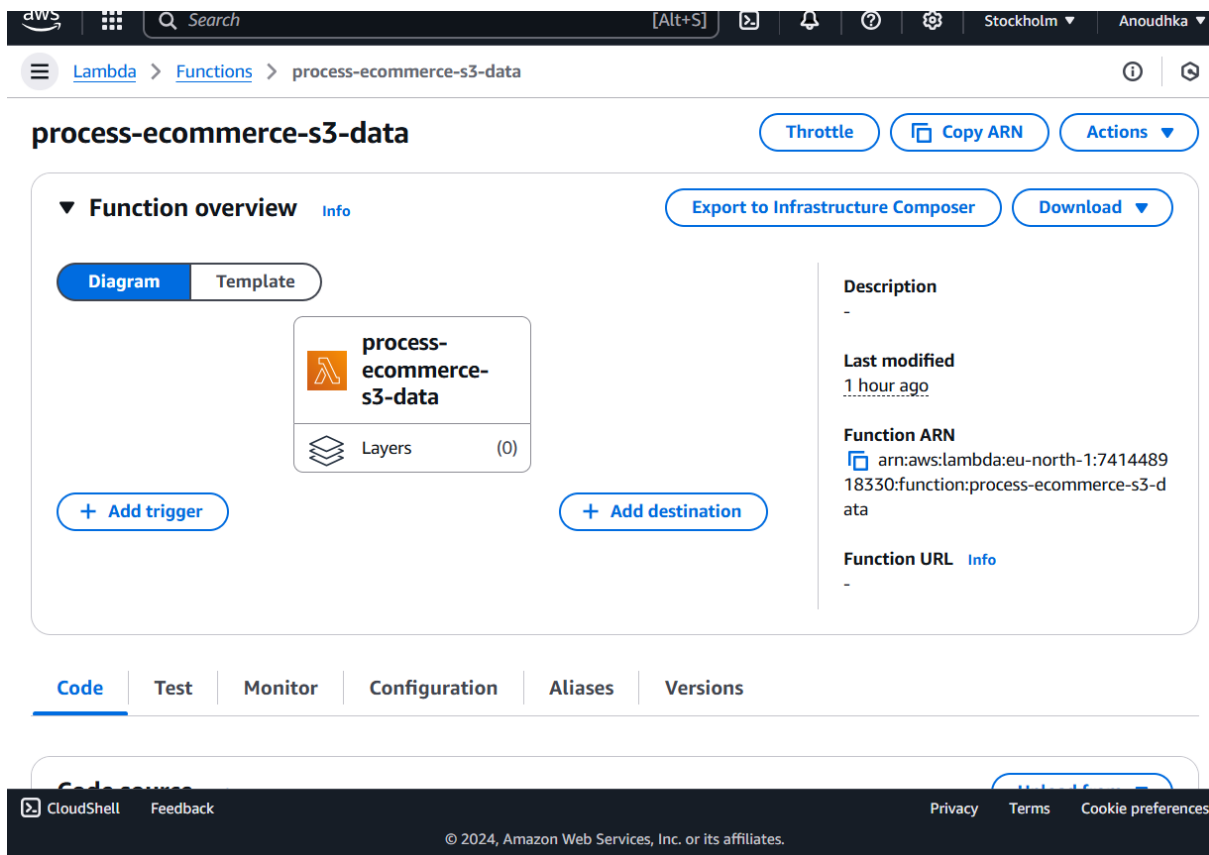
- **Total Products by Sub-Category:**
The "**Topwear**" and "**Blazers, Waistcoats and Suits**" categories dominate the total product count, accounting for a significant portion of the inventory. This indicates a higher focus on these sub-categories within the dataset.
- **Total Money Earned by Brands:**
The **REEB** and **ECKO Unl** brands have the highest total earnings, significantly outpacing other brands. Brands such as **Pu** and **Free Authori** also contribute to overall sales, but their earnings are comparatively lower. This suggests that REEB and ECKO Unl are key revenue drivers.
- **Top 5 Brands by Discount:**
The brands **SIYA SHOPIVA**, **Jack Roy**, and **fas** are offering the highest total discounts. This could indicate a strategy to attract more customers through higher discounts or possibly clear inventory.
- **Product Availability:**
A significant number of products are **in stock** (True), while a smaller portion is marked as **out of stock** (False). This suggests that the inventory is well-maintained, with most products currently available to customers.

BONUS

I also automated the entire pipeline. I started by creating a role that will be accessed by aws lambda to get data from s3 and send out push notifications using sns.



I also created a new function in Lambda. The script for processing can be found in the attachments. I attached the IAM I created to this function.



I then update the permissions of my s3 bucket to include s3 event notification with the trigger destination being the lambda function that I set up.

aws Search [Alt+S] Stockholm Anoudhka

Amazon S3 > Buckets > anoushka-bda

✓ Successfully created event notification "trigger-lambda-on-upload". Operation successfully completed.

Configure CloudTrail data events to log Amazon S3 object-level API operations in the CloudTrail console. [Learn more](#)

Name	Access
No data events	
No data events to display.	

[Configure in CloudTrail](#)

Event notifications (1) [Edit](#) [Delete](#) [Create event notification](#)

Send a notification when specific events occur in your bucket. [Learn more](#)

<input type="checkbox"/>	Name	Event types	Filters	Destination type	Destination
<input type="checkbox"/>	trigger-lambda-on-upload	All object create events	-	Lambda function	process-ecommerce-s3-data

Amazon EventBridge [Edit](#)

For additional capabilities, use Amazon EventBridge to build event-driven applications at scale using S3 event notifications. [Learn more](#) or

Then, I wrote the lambda code that will read the file from S3, process or transform the data (simple operation like format validation) and save processed data back to a different S3 bucket. You can find the code attached in the zip file. After we deploy the lambda function, we will create an SNS topic.

aws Search [Alt+S] Stockholm Anoudhka

Amazon SNS > Topics > data-pipeline-status

Amazon SNS

- Dashboard
- Topics
- Subscriptions
- ▼ Mobile
 - Push notifications
 - Text messaging (SMS)

data-pipeline-status [Edit](#) [Delete](#) [Publish message](#)

Details

Name data-pipeline-status	Display name -
ARN arn:aws:sns:eu-north-1:741448918330:data-pipeline-status	Topic owner 741448918330
Type Standard	

< [Subscriptions](#) [Access policy](#) [Data protection policy](#) [Delivery policy \(HTTP/S\)](#) [Delivery status logging](#) [Encryption](#) [Tags](#) >

Subscriptions (1) [Edit](#) [Delete](#) [Request confirmation](#) [Confirm subscription](#) [Create subscription](#)

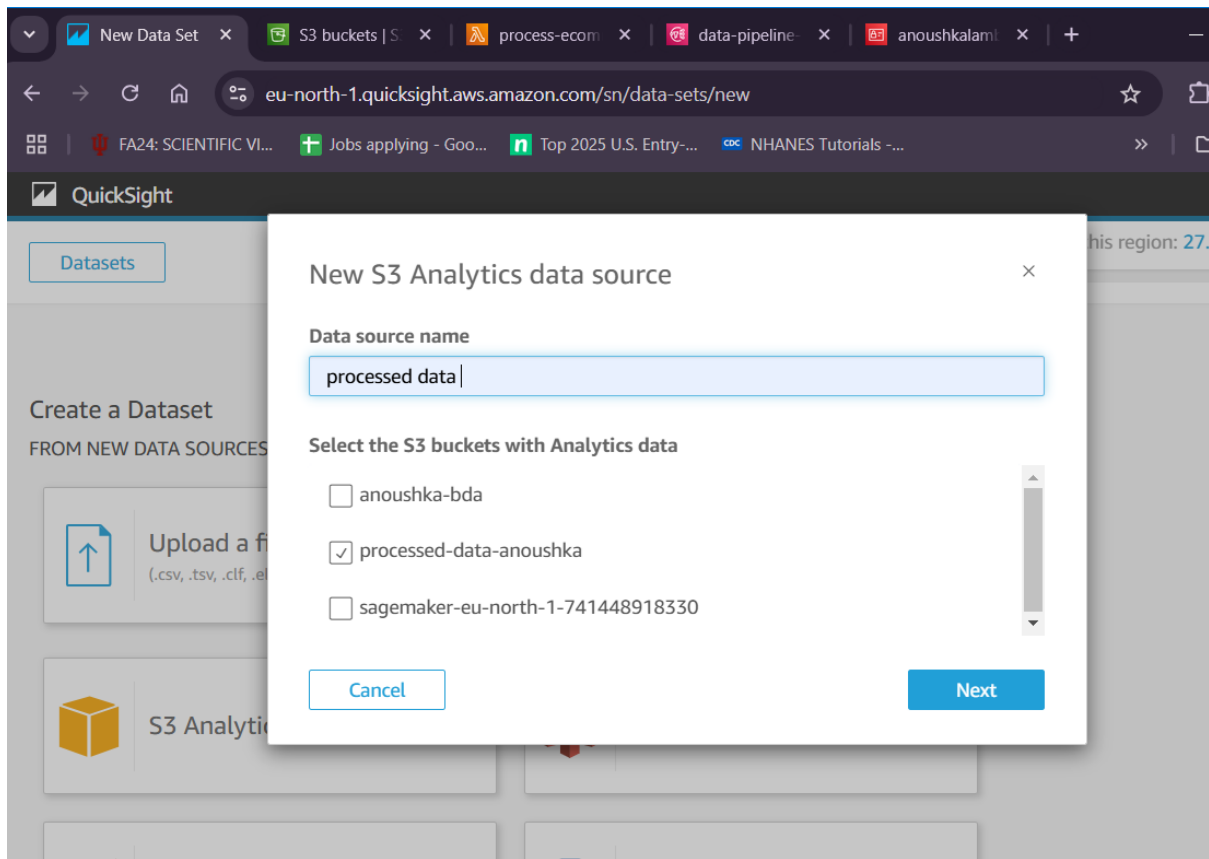
Search

ID	Endpoint	Status	Protocol
<input type="radio"/> bffd38b7-cb6e-4318-8098-eaf9c4...	anshinde@iu.edu	Confirmed	EMAIL

Additionally, I created a new bucket to store all the data coming from the lambda function and linked it to quick insight.

	Name	AWS Region
<input type="radio"/>	anoushka-bda	Europe (Stockholm) eu-north-1
<input type="radio"/>	processed-data-anoushka	Europe (Stockholm) eu-north-1

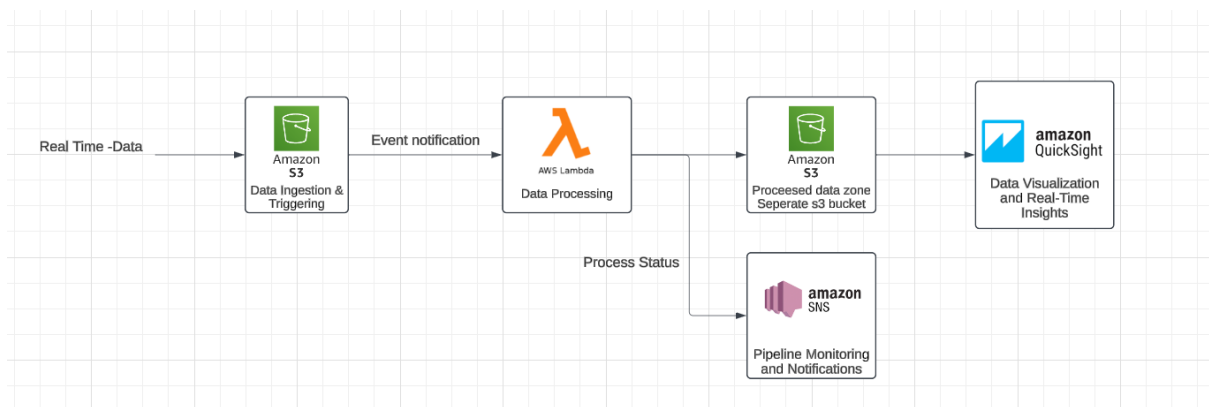
Finally, I connected it to QuickSight for Real Time Insights.



Architecture Summary

1. **Data Ingestion & Triggering:** **S3 Event Notifications** for triggering.
2. **Data Processing:** **AWS Lambda** for lightweight processing.
3. **Data Storage:** **Amazon S3** for raw and processed data.
4. **Notifications:** **AWS SNS** for pipeline status alerts.
5. **Data Visualization:** **Amazon QuickSight** (with SPICE) for dashboards.

Architecture diagram of the pipeline



Results

The implementation of the big data pipeline yielded significant insights into various aspects of the e-commerce platform's operations and performance metrics. The analysis revealed several key findings across multiple dimensions of the business.

The AWS SageMaker Autopilot implementation demonstrated robust predictive capabilities for average rating prediction. The model achieved notable performance metrics:

- Root Mean Square Error (RMSE) of 0.333
- Mean Square Error (MSE) of 0.111

The predicted vs. actual plot shows a strong positive correlation between predicted and actual values, with most predictions falling within the confidence band. The model consistently predicts values within ± 0.333 of the actual average rating, indicating reliable performance for practical applications.

The analysis of product distribution revealed significant insights into inventory composition. A total of approximately 13.4K products were analyzed across various sub-categories. Topwear emerged as the dominant category, representing the largest segment. Tracksuits and Blazers/Waistcoats/Suits formed significant secondary categories. Clothing Accessories and specialized categories like Crocks Club showed moderate representation.

Revenue analysis across brands revealed clear market leaders. REEB emerged as the top-performing brand in terms of revenue generation. ECKO Unltd showed strong performance as the second-highest revenue generator. A clear tiering of brand performance was observed, with significant gaps between top performers and middle-tier brands. Brands like Amp and Keo maintained steady mid-level performance.

The product availability analysis presented crucial operational insights. A substantial portion of the inventory showed as 'False' in the availability metric, suggesting potential stock management challenges. The disparity between available and unavailable products indicates an opportunity for inventory optimization.

Finally, the model's feature importance analysis revealed valuable insights into rating predictions. Seller information emerged as a significant predictor, contributing 5.985% to the model's decisions. Product description and brand showed similar importance levels (approximately 3%). The impact of sellers on rating predictions varied significantly, with some sellers showing consistent positive impact while others demonstrated more variable effects.

Conclusion

This comprehensive analysis of the Flipkart e-commerce dataset through our implemented big data pipeline has revealed several crucial insights for business operations and strategy. The machine learning model's performance demonstrates the viability of automated rating prediction systems, while the business metrics highlight clear opportunities for optimization.

The variation in brand performance and product category distribution suggests opportunities for targeted inventory management and marketing strategies. The significant impact of sellers on product ratings indicates the importance of seller quality management and potential for optimization of seller partnerships.

The implementation successfully demonstrated the scalability and effectiveness of combining AWS services with PySpark for large-scale e-commerce data analysis. The insights generated provide actionable intelligence for business decision-making, particularly in areas of inventory management, brand partnerships, and seller relations.

References:

Kaggle: <https://www.kaggle.com/datasets/aaditshukla/flipkart-fashion-products-dataset/code>

QuickSight Documentation: <https://docs.aws.amazon.com/quicksight/latest/user/quickstart-createanalysis.html>

SageMaker Auto Tutorial: <https://aws.amazon.com/tutorials/machine-learning/tutorial-automatically-create-models/>